

```

1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt

1 train= pd.read_csv('/content/train.csv')
2 test= pd.read_csv('/content/test.csv')
3 sample= pd.read_csv('/content/sample.csv')
4
5 train.info()
6 train.head()

```



<class 'pandas.core.frame.DataFrame'>

RangeIndex: 8693 entries, 0 to 8692

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	PassengerId	8693 non-null	object
1	HomePlanet	8492 non-null	object
2	CryoSleep	8476 non-null	object
3	Cabin	8494 non-null	object
4	Destination	8511 non-null	object
5	Age	8514 non-null	float64
6	VIP	8490 non-null	object
7	RoomService	8512 non-null	float64
8	FoodCourt	8510 non-null	float64
9	ShoppingMall	8485 non-null	float64
10	Spa	8510 non-null	float64
11	VRDeck	8505 non-null	float64
12	Name	8493 non-null	object
13	Transported	8693 non-null	bool

dtypes: bool(1), float64(6), object(7)

memory usage: 891.5+ KB

	PassengerId	HomePlanet	CryoSleep	Cabin	Destination	Age	VIP	RoomService	FoodCourt	S
0	0001_01	Europa	False	B/0/P	TRAPPIST-1e	39.0	False	0.0	0.0	
1	0002_01	Earth	False	F/0/S	TRAPPIST-1e	24.0	False	109.0	9.0	
2	0003_01	Europa	False	A/0/S	TRAPPIST-1e	58.0	True	43.0	3576.0	
3	0003_02	Europa	False	A/0/S	TRAPPIST-1e	33.0	False	0.0	1283.0	
4	0004_01	Earth	False	F/1/S	TRAPPIST-1e	16.0	False	303.0	70.0	

```
1 numerical_columns = train.select_dtypes(include=['float64', 'int64']).columns
2 for i in numerical_columns:
3     train[i].fillna(train[i].mean(), inplace=True)
4
5 categorical_columns = train.select_dtypes(include=['object']).columns
6 for i in categorical_columns:
7     train[i].fillna(train[i].mode()[0], inplace=True)
```

```
1 train.isnull().sum()
```

```
PassengerId      0
HomePlanet        0
CryoSleep         0
Cabin             0
Destination       0
Age               0
VIP               0
RoomService       0
FoodCourt         0
ShoppingMall      0
Spa               0
VRDeck            0
Name              0
Transported       0
dtype: int64
```

```
1 train.columns
```

```
Index(['PassengerId', 'HomePlanet', 'CryoSleep', 'Cabin', 'Destination', 'Age',
      'VIP', 'RoomService', 'FoodCourt', 'ShoppingMall', 'Spa', 'VRDeck',
      'Name', 'Transported'],
      dtype='object')
```

```
1 train.drop(['PassengerId', 'Name', 'Cabin'], axis=1, inplace=True)
2 train
```

	HomePlanet	CryoSleep	Destination	Age	VIP	RoomService	FoodCourt	ShoppingMa
0	Europa	False	TRAPPIST-1e	39.0	False	0.0	0.0	(
1	Earth	False	TRAPPIST-1e	24.0	False	109.0	9.0	28
2	Europa	False	TRAPPIST-1e	58.0	True	43.0	3576.0	(
3	Europa	False	TRAPPIST-1e	33.0	False	0.0	1283.0	371
4	Earth	False	TRAPPIST-1e	16.0	False	303.0	70.0	151
...	...	...	...	...	...	...	...	
8688	Europa	False	55 Cancr i e	41.0	True	0.0	6819.0	(
8689	Earth	True	PSO J318.5-22	18.0	False	0.0	0.0	(
8690	Earth	False	TRAPPIST-1e	26.0	False	0.0	0.0	187%
8691	Europa	False	55 Cancr i e	32.0	False	0.0	1049.0	(
8692	Europa	False	TRAPPIST-1e	44.0	False	126.0	4688.0	(

8693 rows × 11 columns

```
1 X=train.drop('Transported',axis=1)
2 Y=train['Transported']
3 X
```

	HomePlanet	CryoSleep	Destination	Age	VIP	RoomService	FoodCourt	ShoppingMa
0	Europa	False	TRAPPIST-1e	39.0	False	0.0	0.0	(
1	Earth	False	TRAPPIST-1e	24.0	False	109.0	9.0	28
2	Europa	False	TRAPPIST-1e	58.0	True	43.0	3576.0	(
3	Europa	False	TRAPPIST-1e	33.0	False	0.0	1283.0	371
4	Earth	False	TRAPPIST-1e	16.0	False	303.0	70.0	151
...	...	...	...	...	...	...	...	
8688	Europa	False	55 Cancr i e	41.0	True	0.0	6819.0	(
8689	Earth	True	PSO J318.5-22	18.0	False	0.0	0.0	(
8690	Earth	False	TRAPPIST-1e	26.0	False	0.0	0.0	187%
8691	Europa	False	55 Cancr i e	32.0	False	0.0	1049.0	(
8692	Europa	False	TRAPPIST-1e	44.0	False	126.0	4688.0	(

8693 rows × 10 columns

```

1 from sklearn.preprocessing import LabelEncoder
2 # Initialize LabelEncoder
3 label_encoder = LabelEncoder()
4 # Iterate over each column in your dataset
5 for column in X.columns:
6     # Check if the column dtype is 'object' (i.e., categorical)
7     if X[column].dtype == 'object':
8         # Encode the values in the column and replace them with encoded values
9         X[column] = label_encoder.fit_transform(X[column])

1 # This will be needed when we have string values in target column
2
3 # Y_encoder = LabelEncoder()
4 # Y = Y_encoder.fit_transform(train['SalePrice'])

1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(X,Y,test_size=0.2,random_state=1)

```

## ✓ Classification

```

1 from sklearn.metrics import accuracy_score #higher is good
2 from sklearn.metrics import mean_squared_error #lower is good

```

```

1 from sklearn.linear_model import LogisticRegression
2 lr= LogisticRegression()
3 lr.fit(x_train, y_train)
4 y_pred = lr.predict(x_test)
5 # mse = mean_squared_error(y_test, y_pred)
6 # print("Mean Squared Error:", mse)
7 accuracy = accuracy_score(y_test, y_pred)
8 print(f"Accuracy: {accuracy * 100:.2f}%")

```

Accuracy: 79.30%

/usr/local/lib/python3.10/dist-packages/sklearn/linear\_model/\_logistic.py:458: ConvergenceWarning  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

```
1 # Decision Tree Classification
2 from sklearn.tree import DecisionTreeClassifier
3 dt = DecisionTreeClassifier()
4 dt.fit(x_train, y_train)
5 y_pred = dt.predict(x_test)
6 accuracy = accuracy_score(y_test, y_pred)
7 print(f"Accuracy: {accuracy * 100:.2f}%")
```

Accuracy: 74.99%

```
1 from sklearn.ensemble import RandomForestClassifier
2 # rf = RandomForestClassifier(n_estimators=100, random_state=42)
3 rf = RandomForestClassifier()
4 rf.fit(x_train, y_train)
5 y_pred = rf.predict(x_test)
6 accuracy = accuracy_score(y_test, y_pred)
7 print(f"Accuracy: {accuracy * 100:.2f}%")
```

Accuracy: 78.55%

```
1 from sklearn.ensemble import GradientBoostingClassifier
2 gb = GradientBoostingClassifier()
3 gb.fit(x_train,y_train)
4 y_pred=gb.predict(x_test)
5 accuracy = accuracy_score(y_test, y_pred)
6 print(f"Accuracy: {accuracy * 100:.2f}%")
```

Accuracy: 79.70%

```
1 from sklearn.svm import SVC
2 svc=SVC()
3 svc.fit(x_train,y_train)
4 y_pred=svc.predict(x_test)
5 accuracy = accuracy_score(y_test, y_pred)
6 print(f"Accuracy: {accuracy * 100:.2f}%")
```

Accuracy: 79.93%

```
1 from xgboost import XGBClassifier
2 xgb = XGBClassifier()
3 xgb.fit(x_train,y_train)
4 y_pred=xgb.predict(x_test)
5 accuracy = accuracy_score(y_test, y_pred)
6 print(f"Accuracy: {accuracy * 100:.2f}%")
```

Accuracy: 78.95%

```

1 from xgboost import XGBRFClassifier
2 xgbclf = XGBRFClassifier()
3 xgbclf.fit(x_train,y_train)
4 y_pred=xgbclf.predict(x_test)
5 accuracy = accuracy_score(y_test, y_pred)
6 print(f"Accuracy: {accuracy * 100:.2f}%")

```

Accuracy: 79.70%

```

1 from lightgbm import LGBMClassifier
2 lgb = LGBMClassifier()
3 lgb.fit(x_train,y_train)
4 y_pred=lgb.predict(x_test)
5 accuracy = accuracy_score(y_test, y_pred)
6 print(f"Accuracy: {accuracy * 100:.2f}%")

```

```

[LightGBM] [Info] Number of positive: 3482, number of negative: 3472
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.001320
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1363
[LightGBM] [Info] Number of data points in the train set: 6954, number of used features: 10
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500719 -> initscore=0.002876
[LightGBM] [Info] Start training from score 0.002876
Accuracy: 79.36%

```

```

1 from sklearn.neural_network import MLPClassifier
2 # mlp = MLPClassifier(hidden_layer_sizes=(100,), max_iter=500)
3 mlp = MLPClassifier()
4 mlp.fit(x_train,y_train)
5 y_pred=mlp.predict(x_test)
6 accuracy = accuracy_score(y_test, y_pred)
7 print(f"Accuracy: {accuracy * 100:.2f}%")

```

Accuracy: 78.61%

```

1 from sklearn.neighbors import KNeighborsClassifier
2 knn = KNeighborsClassifier(n_neighbors=5)
3 knn.fit(x_train, y_train)
4 y_pred = knn.predict(x_test)
5 accuracy = accuracy_score(y_test, y_pred)
6 print(f"Accuracy: {accuracy * 100:.2f}%")

```

Accuracy: 78.32%

```
1 from sklearn.naive_bayes import GaussianNB
2 nb = GaussianNB()
3 nb.fit(x_train, y_train)
4 y_pred = nb.predict(x_test)
5 accuracy = accuracy_score(y_test, y_pred)
6 print(f"Accuracy: {accuracy * 100:.2f}%")
```

Accuracy: 70.85%

```
1 from sklearn.ensemble import AdaBoostClassifier
2 ada = AdaBoostClassifier(n_estimators=100)
3 ada.fit(x_train, y_train)
4 y_pred = ada.predict(x_test)
5 accuracy = accuracy_score(y_test, y_pred)
6 print(f"Accuracy: {accuracy * 100:.2f}%")
```

Accuracy: 79.82%

```
1 t=test
```

```
1 numerical_columns = t.select_dtypes(include=['float64', 'int64']).columns
2 for i in numerical_columns:
3     t[i].fillna(t[i].mean(), inplace=True)
4
5 categorical_columns = t.select_dtypes(include=['object']).columns
6 for i in categorical_columns:
7     t[i].fillna(t[i].mode()[0], inplace=True)
```

```
1 t.isnull().sum()
```

```
PassengerId    0
HomePlanet     0
CryoSleep      0
Cabin          0
Destination    0
Age            0
VIP            0
RoomService    0
FoodCourt      0
ShoppingMall   0
Spa            0
VRDeck         0
Name           0
dtype: int64
```

```
1 t.drop(['PassengerId', 'Name', 'Cabin'], axis=1, inplace=True)
2 t
```

	HomePlanet	CryoSleep	Destination	Age	VIP	RoomService	FoodCourt	Shopp
0	Earth	True	TRAPPIST-1e	27.000000	False	0.0	0.0	
1	Earth	False	TRAPPIST-1e	19.000000	False	0.0	9.0	
2	Europa	True	55 Cancr i e	31.000000	False	0.0	0.0	
3	Europa	False	TRAPPIST-1e	38.000000	False	0.0	6652.0	
4	Earth	False	TRAPPIST-1e	20.000000	False	10.0	0.0	
...	...	...	...	...	...	...	...	...
4272	Earth	True	TRAPPIST-1e	34.000000	False	0.0	0.0	
4273	Earth	False	TRAPPIST-1e	42.000000	False	0.0	847.0	
4274	Mars	True	55 Cancr i e	28.658146	False	0.0	0.0	
4275	Europa	False	TRAPPIST-1e	28.658146	False	0.0	2680.0	
4276	Earth	True	PSO J318.5-22	43.000000	False	0.0	0.0	

4277 rows × 10 columns

```

1 from sklearn.preprocessing import LabelEncoder
2
3 # Initialize LabelEncoder
4 label_encoder = LabelEncoder()
5
6 # Iterate over each column in your dataset
7 for column in t.columns:
8     # Check if the column dtype is 'object' (i.e., categorical)
9     if t[column].dtype == 'object':
10         # Encode the values in the column and replace them with encoded values
11         t[column] = label_encoder.fit_transform(t[column])

```

```

1 svc.fit(X,Y)
2 y_pred=svc.predict(test)

```

```

1 y_pred

array([ True, False,  True, ...,  True,  True,  True])

```

```

1 # it will bw only used when we have label encoded object Target column
2
3 # y_pred_decoded = Y_encoder.inverse_transform(y_pred)

```

```

1 sample1=sample

```



```
1 sample1['Transported']=y_pred
```

```
1 sample1
```

	PassengerId	Transported
0	0013_01	True
1	0018_01	False
2	0019_01	True
3	0021_01	True
4	0023_01	True
...	...	...
4272	9266_02	True
4273	9269_01	True
4274	9271_01	True
4275	9273_01	True
4276	9277_01	True

4277 rows × 2 columns

```
1 sample1.to_csv('submit_1.csv',index=False)
```

## ✓ ANN

```
1 x_train = x_train.astype(np.int32)
2 y_train = y_train.astype(np.int32)
3 x_test = x_test.astype(np.int32)
4 y_test = y_test.astype(np.int32)
```

```
1 import tensorflow
2 from tensorflow import keras
3 from tensorflow.keras import Sequential
4 from tensorflow.keras.layers import Dense
```

```
1 model=Sequential()
```

```

1 model.add(Dense(1000, activation='relu',input_dim=10))
2 model.add(Dense(1000, activation='relu'))
3 model.add(Dense(1000, activation='relu'))
4 model.add(Dense(500, activation='relu'))
5 model.add(Dense(500, activation='relu'))
6 model.add(Dense(100, activation='relu'))
7 model.add(Dense(1, activation='relu'))

```

```
1 model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
dense_14 (Dense)	(None, 1000)	11000
dense_15 (Dense)	(None, 1000)	1001000
dense_16 (Dense)	(None, 1000)	1001000
dense_17 (Dense)	(None, 500)	500500
dense_18 (Dense)	(None, 500)	250500
dense_19 (Dense)	(None, 100)	50100
dense_20 (Dense)	(None, 1)	101
=====		
Total params: 2814201 (10.74 MB)		
Trainable params: 2814201 (10.74 MB)		
Non-trainable params: 0 (0.00 Byte)		
=====		

```
1 model.compile(loss='binary_crossentropy',optimizer='Adam',metrics=['accuracy'])
```

```
1 model.fit(x_train, y_train, batch_size=200, validation_data=(x_test, y_test), epochs=
```

```

Epoch 1/10
35/35 [=====] - 3s 96ms/step - loss: 7.7236 - accuracy: 0.4993 - val_1
Epoch 2/10
35/35 [=====] - 5s 138ms/step - loss: 7.7236 - accuracy: 0.4993 - val_
Epoch 3/10
35/35 [=====] - 3s 85ms/step - loss: 7.7236 - accuracy: 0.4993 - val_1
Epoch 4/10
35/35 [=====] - 3s 95ms/step - loss: 7.7236 - accuracy: 0.4993 - val_1
Epoch 5/10
35/35 [=====] - 3s 90ms/step - loss: 7.7236 - accuracy: 0.4993 - val_1
Epoch 6/10
35/35 [=====] - 4s 129ms/step - loss: 7.7236 - accuracy: 0.4993 - val_
Epoch 7/10
35/35 [=====] - 3s 86ms/step - loss: 7.7236 - accuracy: 0.4993 - val_1
Epoch 8/10
35/35 [=====] - 4s 107ms/step - loss: 7.7236 - accuracy: 0.4993 - val_
Epoch 9/10
35/35 [=====] - 3s 96ms/step - loss: 7.7236 - accuracy: 0.4993 - val_1

```

Epoch 10/10

35/35 [=====] - 4s 128ms/step - loss: 7.7236 - accuracy: 0.4993 - val\_<keras.src.callbacks.History at 0x7b8f5a3e8550>

```
1 X =X.astype(np.int32)
2 Y = Y.astype(np.int32)
3 model.fit(X,Y,epochs=10)
```

Epoch 1/10

272/272 [=====] - 11s 40ms/step - loss: 7.7684 - accuracy: 0.4964

Epoch 2/10

272/272 [=====] - 12s 43ms/step - loss: 7.7684 - accuracy: 0.4964

Epoch 3/10

272/272 [=====] - 12s 43ms/step - loss: 7.7684 - accuracy: 0.4964

Epoch 4/10

272/272 [=====] - 12s 43ms/step - loss: 7.7684 - accuracy: 0.4964

Epoch 5/10

272/272 [=====] - 13s 49ms/step - loss: 7.7684 - accuracy: 0.4964

Epoch 6/10

272/272 [=====] - 12s 43ms/step - loss: 7.7684 - accuracy: 0.4964

Epoch 7/10

272/272 [=====] - 17s 62ms/step - loss: 7.7684 - accuracy: 0.4964

Epoch 8/10

272/272 [=====] - 14s 50ms/step - loss: 7.7684 - accuracy: 0.4964

Epoch 9/10

272/272 [=====] - 20s 73ms/step - loss: 7.7684 - accuracy: 0.4964

Epoch 10/10

272/272 [=====] - 19s 69ms/step - loss: 7.7684 - accuracy: 0.4964  
<keras.src.callbacks.History at 0x7b8f480e2bc0>

```
1 test = test.astype(np.int32)
2 y_pred=model.predict(test)
```

134/134 [=====] - 1s 8ms/step

```
1 sample2 = sample
```

```
1 sample2['Transported']=y_pred
```

```
1 sample2
```